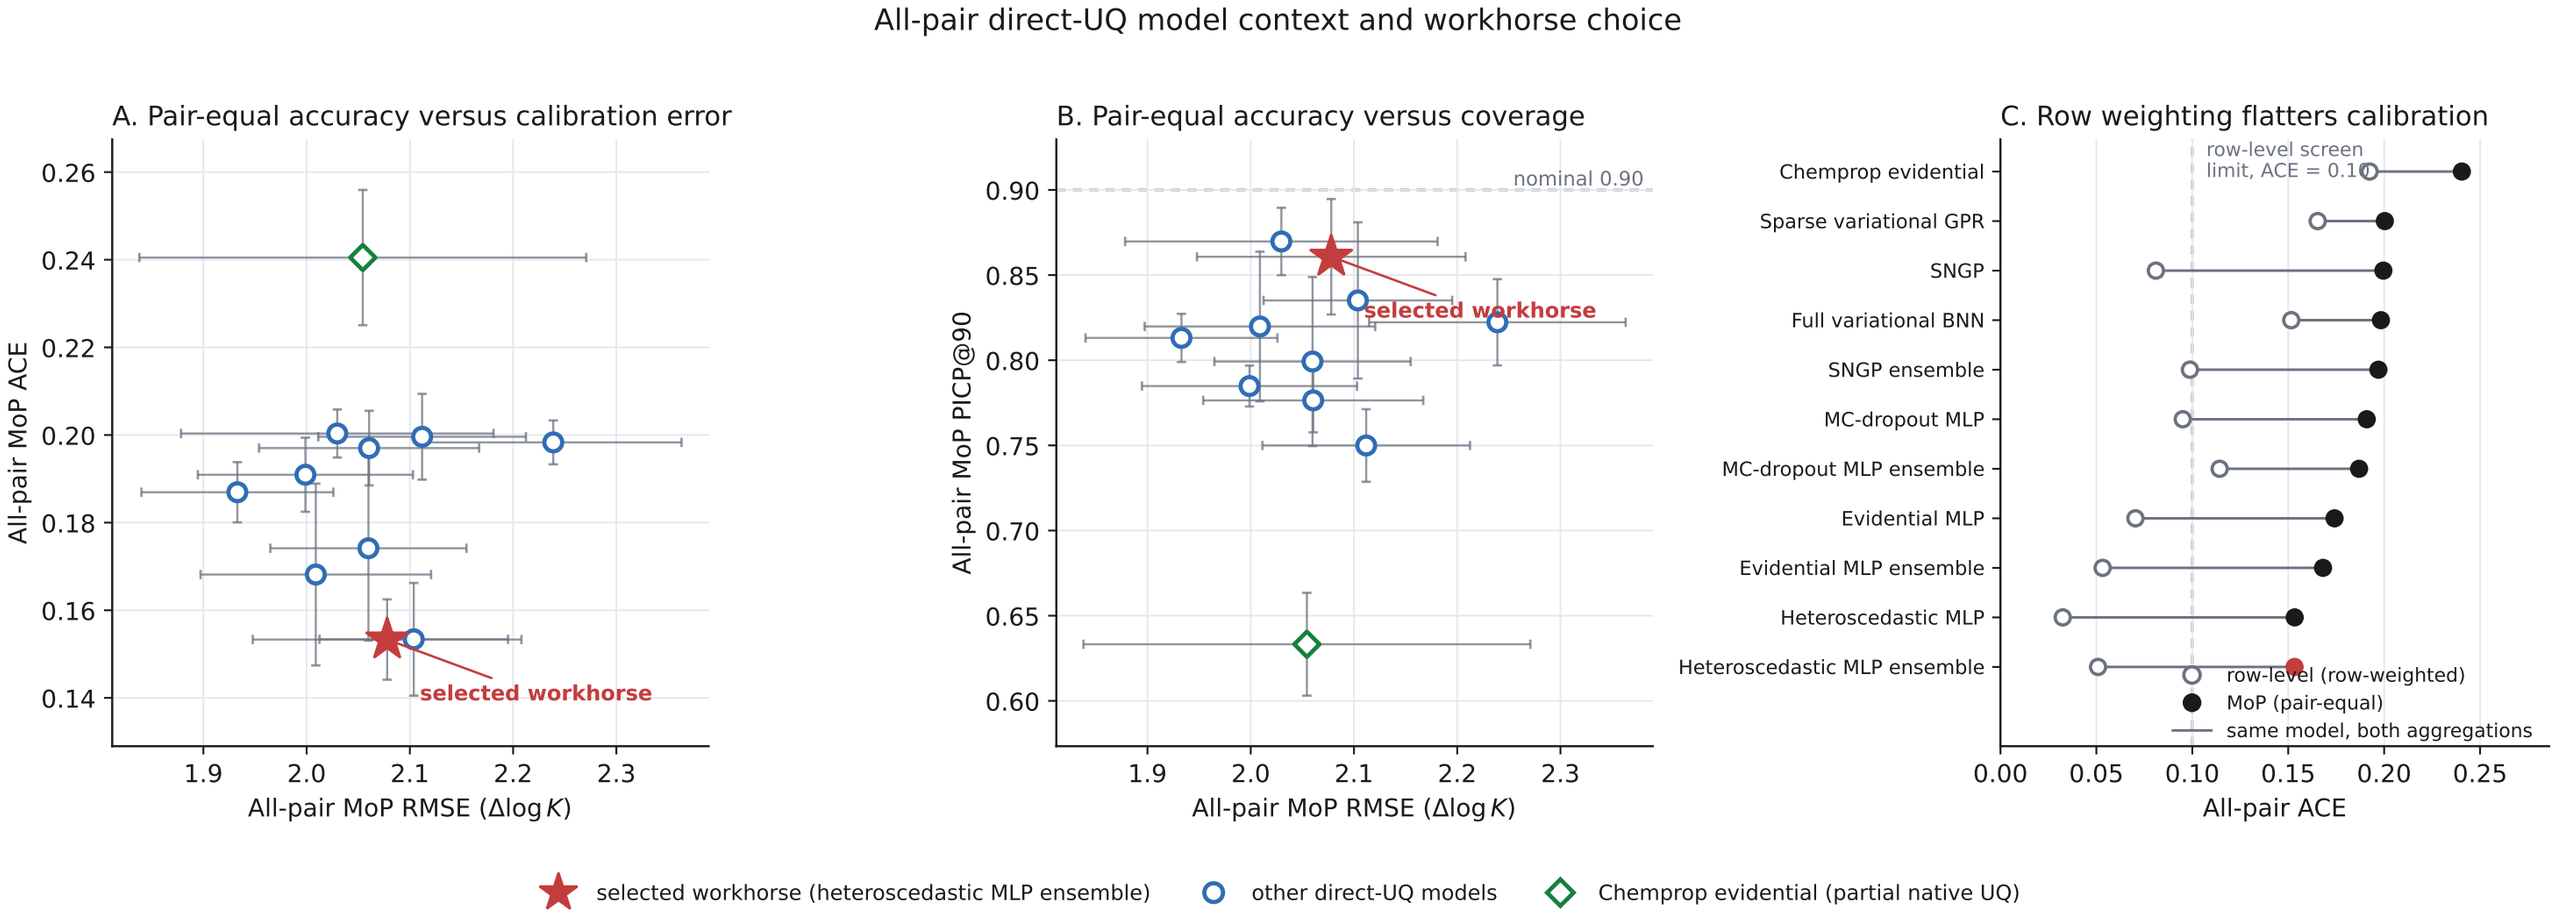

Panels A and B are pair-equal (MoP) throughout; error bars are fold SD ($n = 5$, ddof = 1). Accuracy differences across the grid are smaller than those bars, so the candidates separate on calibration rather than on RMSE. Panel C gives the row-weighted value of the same quantity: every model's ACE is lower row-weighted than pair-equal (median ratio 2.0 ×), because row weighting is dominated by densely sampled pairs. Row-level values are reported only here.