Objective-specific pair acquisition

A. High selectivity ($\Delta\log K > 1$)

higher is better

True hit@k

1.0
0.9
0.8
0.7
0.6

at $k \geq 5$ th trongest penalty
finds an event ss often than random

$k = 1$ $k = 5$ $k = 10$

B. Near-zero selectivity ($|\Delta\log K| < 1$)

higher is better

1.0
0.9
0.8
0.7
0.6

$k = 1$ $k = 5$ $k = 10$

C. High selectivity ($\Delta\log K > 1$)

higher is better

threshold margin leads
where LCB collapses

Event precision@k

0.9
0.8
0.7
0.6
0.5
0.4

$k = 1$ $k = 5$ $k = 10$

D. Near-zero selectivity ($|\Delta\log K| < 1$)

higher is better

and matches the best rule here

0.9
0.8
0.7
0.6
0.5
0.4

$k = 1$ $k = 5$ $k = 10$

E. High selectivity ($\Delta\log K > 1$)

lower is better ↓

Retrospective regret@k ($\Delta\log K$)

4
3
2
1
0

$k = 1$ $k = 5$ $k = 10$

F. Near-zero selectivity ($|\Delta\log K| < 1$)

lower is better ↓

1.75
1.50
1.25
1.00
0.75
0.50
0.25
0.00

$k = 1$ $k = 5$ $k = 10$

predicted selectivity LCB-1 LCB-2 LCB-3 threshold margin random

LCB-1/2/3 are one rule at increasing uncertainty penalty, shaded light to dark. The threshold-margin rule divides the distance past the event threshold by σ_{total} instead of subtracting it, and is the only uncertainty-integrated rule that never falls materially below the value-only baseline on an event metric in either objective.
Bars are fold means, error bars fold SD ($n = 5$), after a pair-equal mean within each fold.